

Tanay Mehendale

San Jose, CA (can relocate) | tanay.mehendale@tamu.edu | 979-344-3679 | [LinkedIn/tanay-mehendale/](https://www.linkedin.com/in/tanay-mehendale/) | [Portfolio](#)

SUMMARY

Data Engineer with 2+ years of hands-on experience in building ETL pipelines, automating data workflows, and bridging SAP systems with cloud platforms. Proficient in both traditional enterprise data (SAP, SQL Server) and modern tools like Spark, Airflow, and AWS. Currently learning Snowflake and GenAI to modernize data workflows.

EXPERIENCE

Data Engineer (SAP Technical Consultant) - LTIMindtree | Mumbai, India Jan 2021 – Jun 2023

- Developed SAP reports and scripts using **SQL** and ABAP to extract, filter, and join transactional data, reducing upstream data mismatches by 18%
- Built an **ETL** data pipeline using **Spark** to integrate structured data from 7 data sources into **AWS Redshift**, processing over 750 GB/month data, and reducing reporting latency for operational KPIs
- Created modular **Python scripts** to clean and validate SAP export files post-extraction, enhancing automated testing and reducing ingestion errors across 50+ weekly batch runs
- Managed **orchestration** and monitoring of 30+ Spark jobs in **Airflow**, configuring DAG schedules, retries and SLA windows to improve pipeline reliability and reduce ingestion delays by 12%
- Supported **Git** based **version control** and peer review process for SAP extract logic and ETL scripts to maintain code consistency across deployment cycles
- Documented 25+ data workflows for new developers to easily understand existing operational workflow, reduce onboarding time and create an effective knowledge base
- Led **cross functional collaboration** with global teams and the leadership across 52+ agile product releases, simplifying SDLC audit process to eliminate compliance breaches and reduce release issues by 84%

Data Analyst - Texas A&M University | College Station, TX Apr 2024 – May 2025

- Automated **data ingestion** using scripts for 54 raw files resulting in a 60% reduction in manual effort
- Analyzed student compliance records in **Excel** using pivot tables and VLOOKUP to spot recurring submission errors across 5 departments, helping reduce exception cases by 20% over one semester
- **Mentored** 25+ students on data modeling best practices and translating business needs into structured datasets during Business Systems Analysis coursework

EDUCATION

Texas A&M University - M.S. Management Information Systems | College Station, TX May 2025

University of Mumbai - B.Tech. Electronics & Communications | Mumbai, India May 2021

PROJECTS

Real Time Sign-Up Health Monitor | [Link](#) | Airflow, Kafka, Spark, Cassandra, Docker, Git Aug 2025

- Built a real-time sign-up health pipeline that detects major drops (>50%) within 2 minutes in local tests, resulting in 90% faster detection than 60min batch checks

Data Warehouse for Retail Analytics | [Link](#) | Microsoft SQL Server, SSIS, Tableau Dec 2024

- Designed a retail data warehouse using Kimball approach to analyze \$600M+ retail sales patterns across 2 data marts and 6 years of data, enabling inventory decisions and retail trend analysis through SSRS and Tableau

Point of Sale (POS) System | [Link](#) | AWS, Data Modeling, Clustering & Replication Nov 2023

- Built a POS system on AWS EC2 with relational database MariaDB, implementing clustering and sharding techniques; Optimized read performance by 45% by performing NoSQL data modeling to restructure data

Customer Churn Analysis | [Link](#) | Tableau, Data Visualization, Business Intelligence Jul 2024

- Developed interactive Tableau dashboards analyzing churn patterns and customer segmentation across 6,600+ telecom accounts, supporting targeted retention strategies.

SKILLS

Certifications: [Assoc. Data Engineer](#); *Ongoing* - Snowflake SnowPro Core, AWS Solutions Architect Associate

Coding: SQL, Python, PostgreSQL, Git, R; **Database:** SAP HANA, MariaDB, MongoDB; *Familiar* – Cassandra

Data Engineering: ETL/ELT (Extract, Transform, Load), Microsoft SQL Server; *Familiar* – Shell, Bash

Cloud: AWS EC2, Redshift, S3, Apache Airflow, Snowflake; *Familiar* – Apache Spark, Kafka

BI Tools: Tableau, SSRS (SQL Server Reporting Services); *Familiar* – Power BI

AI (familiar): Prompt Engineering, Snowflake Cortex AI, LLMs, RAG, Cursor, GitHub Copilot